

Create speech



POST <https://api.openai.com/v1/audio/speech>

Generates audio from the input text.

Request body

input string **Required**

The text to generate audio for. The maximum length is 4096 characters.

model string **Required**

One of the available [TTS models](#): `tts-1` , `tts-1-hd` or `gpt-4o-mini-tts` .

voice string **Required**

The voice to use when generating the audio. Supported voices are `alloy` , `ash` , `ballad` , `coral` , `echo` , `fable` , `onyx` , `nova` , `sage` , `shimmer` , and `verse` . Previews of the voices are available in the [Text to speech guide](#).

instructions string Optional

Control the voice of your generated audio with additional instructions. Does not work with `tts-1` or `tts-1-hd` .

response_format string Optional Defaults to mp3

The format to audio in. Supported formats are `mp3` , `opus` , `aac` , `flac` , `wav` , and `pcm` .

speed number Optional Defaults to 1

The speed of the generated audio. Select a value from `0.25` to `4.0` . `1.0` is the default.

stream_format string Optional Defaults to audio

The format to stream the audio in. Supported formats are `sse` and `audio` . `sse` is not supported for `tts-1` or `tts-1-hd` .

Returns

The audio file content or a [stream of audio events](#).

Default

SSE Stream Format

Example request

```
curl https://api.openai.com/v1/audio/speech \
  -H "Authorization: Bearer $OPENAI_API_KEY" \
  -H "Content-Type: application/json" \
```

```
-d '{
  "model": "gpt-4o-mini-tts",
  "input": "The quick brown fox jumped over the lazy dog.",
  "voice": "alloy"
}' \
--output speech.mp3
```

Create transcription



POST <https://api.openai.com/v1/audio/transcriptions>

Transcribes audio into the input language.

Request body

file file **Required**

The audio file object (not file name) to transcribe, in one of these formats: flac, mp3, mp4, mpeg, mpga, m4a, ogg, wav, or webm.

model string **Required**

ID of the model to use. The options are `gpt-4o-transcribe` , `gpt-4o-mini-transcribe` , and `whisper-1` (which is powered by our open source Whisper V2 model).

chunking_strategy "auto" or object **Optional**

Controls how the audio is cut into chunks. When set to `"auto"` , the server first normalizes loudness and then uses voice activity detection (VAD) to choose boundaries. `server_vad` object can be provided to tweak VAD detection parameters manually. If unset, the audio is transcribed as a single block.

▼ Show possible types

include[] array **Optional**

Additional information to include in the transcription response. `logprobs` will return the log probabilities of the tokens in the response to understand the model's confidence in the transcription. `logprobs` only works with `response_format` set to `json` and only with the models `gpt-4o-transcribe` and `gpt-4o-mini-transcribe` .

language string **Optional**

The language of the input audio. Supplying the input language in ISO-639-1 (e.g. `en`) format will improve accuracy and latency.

prompt string **Optional**

An optional text to guide the model's style or continue a previous audio segment. The prompt should match the audio language.

response_format string **Optional** Defaults to json

The format of the output, in one of these options: `json` , `text` , `srt` , `verbose_json` , or `vtt` . For `gpt-4o-transcribe` and `gpt-4o-mini-transcribe` , the only supported format is `json` .

stream `boolean` Optional Defaults to `false`

If set to `true`, the model response data will be streamed to the client as it is generated using [server-sent events](#). See the [Streaming](#) section of the [Speech-to-Text](#) guide for more information.

Note: Streaming is not supported for the `whisper-1` model and will be ignored.

temperature `number` Optional Defaults to `0`

The sampling temperature, between `0` and `1`. Higher values like `0.8` will make the output more random, while lower values like `0.2` will make it more focused and deterministic. If set to `0`, the model will use [log probability](#) to automatically increase the temperature until certain thresholds are hit.

timestamp_granularities[] `array` Optional Defaults to `segment`

The timestamp granularities to populate for this transcription. `response_format` must be set `verbose_json` to use timestamp granularities. Either or both of these options are supported: `word` , or `segment` . Note: There is no additional latency for segment timestamps, but generating word timestamps incurs additional latency.

Returns

The [transcription object](#), a [verbose transcription object](#) or a [stream of transcript events](#).

- Default
- Streaming
- Logprobs
- Word timestamps
- Segment timestamps

Example request

```
curl https://api.openai.com/v1/audio/transcriptions \
-H "Authorization: Bearer $OPENAI_API_KEY" \
-H "Content-Type: multipart/form-data" \
-F file="@/path/to/file/audio.mp3" \
-F model="gpt-4o-transcribe"
```

Response

```
{
  "text": "Imagine the wildest idea that you've ever had, and you're curious about how it might scale 1",
  "usage": {
    "type": "tokens",
    "input_tokens": 14,
    "input_token_details": {
      "text_tokens": 0,
      "audio_tokens": 14
    },
    "output_tokens": 45,
    "total_tokens": 59
  }
}
```

Create translation



POST <https://api.openai.com/v1/audio/translations>

Translates audio into English.

Request body

file file **Required**

The audio file object (not file name) translate, in one of these formats: flac, mp3, mp4, mpeg, mpga, m4a, ogg, wav, or webm.

model string or "whisper-1" **Required**

ID of the model to use. Only `whisper-1` (which is powered by our open source Whisper V2 model) is currently available.

prompt string **Optional**

An optional text to guide the model's style or continue a previous audio segment. The prompt should be in English.

response_format string **Optional** Defaults to json

The format of the output, in one of these options: `json` , `text` , `srt` , `verbose_json` , or `vtt` .

temperature number **Optional** Defaults to 0

The sampling temperature, between 0 and 1. Higher values like 0.8 will make the output more random, while lower values like 0.2 will make it more focused and deterministic. If set to 0, the model will use log probability to automatically increase the temperature until certain thresholds are hit.

Returns

The translated text.

Example request

```
curl https://api.openai.com/v1/audio/translations \
  -H "Authorization: Bearer $OPENAI_API_KEY" \
  -H "Content-Type: multipart/form-data" \
  -F file="@/path/to/file/german.m4a" \
  -F model="whisper-1"
```

Response

```
{  
  "text": "Hello, my name is Wolfgang and I come from Germany. Where are you heading today?"  
}
```
